

Autobelay

long premium, short leash



An auto belay catches a climber with nobody holding the other end. Here the brake is code.

William P. McCormick / William C. McCormick · account [PA3V539Y5LE2](#) · MIT

Alpaca AI Trading Agents Hackathon · Aug 28 – Sep 4, 2026

Autobelay **long premium, short leash**

The thesis, in one sentence

Buy defined-risk, short-dated options premium when an open model can **name a reason**, and let deterministic code size it, stop it, and close it.



● WHY LONG PREMIUM


Long premium only, so the worst case is known before entry, which is what lets every guardrail be a flat number.

Max loss is the debit paid. Nothing can be assigned against us.

The division of labour: the model reads the market and proposes; deterministic code sizes, stops and closes, and is the only thing that can reach an order. The next slide names every piece.

And the autonomy: the trading is autonomous — no human approves an order. The risk *envelope* is not. The whitelist, the position caps and the 2% daily-loss cutoff are git-only, and every knob that *is* reachable at runtime expires at the close.


One cycle, every 10 minutes



The model gets one box. Its proposals and code's own exit sells pass the same gate.

Greeks: Alpaca's, ours as the backstop

We spent three days believing Alpaca's free feed carried no Greeks. We were wrong.



● FROM ALPACA


5,721 of 6,114 contracts in today's chain: one vendor surface, so a strike's put and call deltas agree.

● BLACK-SCHOLES BACKSTOP

Our bisection solve covers the $\sim 6\%$ too thin to price. Every contract records which of the two priced it.

What the model may choose from

One API page is 500 contracts. On SPY, that is three days.



● BEFORE

500 contracts, one page, out to 3 DTE. NVDA's twelve were all one strike.


● NOW

2,912 contracts, three pages, out to 45. At-the-money *and* ~ 0.40 delta, three expiries, both sides, journaled every cycle.

A second opinion

Two crowds answer the same question: Kalshi's index-close market and the option chain's own implied odds. Both are priors to weigh rather than signals to copy, and no code acts on either.

a thin market still quotes every bucket



shown
volume 756



withheld
volume 45

● KNOWING WHEN TO IGNORE IT


A barely-traded market quotes every bucket just as confidently, and nothing about its numbers looks broken. That is the point.

Volume and entropy gates withheld QQQ, **with the reason**, so "no second opinion today" is answerable rather than absent.

...and then we grade it

A prior nobody scores is decoration. Every night, the probabilities the model was handed are Brier-scored against what the market actually did - mean squared error of a probability forecast, where **0 is perfect and 0.25 is a coin flip**.

Brier score - lower is better



● TUE SEP 1, JUDGED ACCOUNT

Kalshi **0.004**, the chain **0.008**. Priors the gate withheld score 0.006, shadow-graded, so we know it discarded nothing good.

One day, and a directional one. What is graded here is the *input* the model was handed, not the trade it made.

...and we grade what it says

The reason strings are load-bearing — the digest, the reviewer, the email, the viewer and this deck all quote them. So every figure the model cites is checked against the prior it was actually shown, and every exit reason against the account. Journaled per cycle, surfaced in the digest, and **reporting only**: prose is not an order parameter, and `check_order` does not grade rhetoric.

`unsupported`

a probability in a prior-shaped clause that matches nothing either crowd handed it

`misattributed`

a real number under the wrong ticker — QQQ's chain prior was withheld, and the model quoted SPY's as "the options market" for QQQ

`fabricated_urgency`

"forced expiry sale", "backstop forces exit" — on a 7-DTE put no code path would touch for six more days

`wrong_direction`

"market held above prior close" with SPY 0.73% *below* it: the strike, 760, read as the prior close, 766.87

`wrong_reference`

a prior close stated outright that is not the prior close

Tue Sep 1, judged account: **22 quoted figures, 22 supported** — and four exit attempts on one losing put whose every stated reason was false. Both findings came out of the journal, not a hunch.

It investigates before it acts

```
research: get_bars({'symbol': 'SPY', 'timeframe': '15Min'})    -> 3334 chars
research: get_news({'symbols': 'NVDA', 'limit': 5})            -> 1770 chars
model output: [{"symbol":"NVDA260902P00220000","side":"buy","qty":10,
  "reason":"NVDA in a clear intraday downtrend (225 -> 218.6); the 220 put
    with 4 DTE, delta -0.58, aligns with the bearish momentum"}]
```

It may call four read-only tools, six times at most, and every call is journaled.
Nothing that places an order is reachable.


The leash

● PER PROPOSAL


whitelist · $\leq \$5,000$ per position · ≤ 4 positions · ≤ 10 contracts
· 2–45 DTE on entries · entries 09:45–15:15 ET · a resting
buy counts as held

● PER DAY

2% loss $\rightarrow$ flatten and halt. A `HALT` file is a global kill switch.



Everything is journaled



Watching it think, from anywhere

The journal streams to a browser at **bot.wpmccormick.pw**, published through a Cloudflare tunnel. There is no VPN and nothing to install.

```
14:30:17 official 7 ORDER Sell 10 SPY200900700/0000 @ 4.0055
Existing 8-DTE put at 760 strike is underwater (entry 4.57, current 3.92) with SPY at 761.28; thesis failed as market held above prior close despite prediction-market bearishness, and
forced close approaches with 7 DTE remaining-cut loss before theta decay accelerates and backstop forces exit.
14:30:17 official ◀ end, 1 action(s)
14:40:09 mixed ♦ PRIOR SPY ref 7686.14 median 7637.5 (-0.63%) P(above) 0.069 vol 11727 → shown
♦ PRIOR QQQ ref 29456.97 median 29050 (-1.38%) P(above) 0.058 vol 614.1 → shown
14:40:09 official ♦ PRIOR SPY ref 7686.14 median 7637.5 (-0.63%) P(above) 0.069 vol 11727 → shown
♦ PRIOR QQQ ref 29456.97 median 29050 (-1.38%) P(above) 0.058 vol 614.1 → shown
14:40:10 mixed ▶ CYCLE equity $100,151.92 day P&L 968.99 positions 0
14:40:10 official ▶ CYCLE equity $100,209.83 day P&L 661.60 positions 0
```

The viewer found four bugs in two days, and none was a failing test: in each case the code did exactly what it had been told.

● #174 #170

The feed cut the model's reasoning mid-word. Closing a +\$1,340 winner, it named a *neighbouring strike*, three cycles running.

● #171 #173

An unfilled limit sat invisible and the same thesis was bought again next door. A ten-minute-old position was proposed for exit, and every number it cited had moved *in its favour*.

Trading

Official

Test / challenge

🎮 Official — trading

🎮 Test / challenge

👤 State

👤 State

👁️ Equity 99,552.71 USD

👁️ Equity 99,552.71 USD

📈 Day P&L -655.10 USD

📈 Day P&L -655.10 USD

👛 Open positions 2

👛 Open positions 2

⬛ Halt state none

⬛ Halt state none


👤 Last decision hold

👤 Last decision hold

Day P&L

Day P&L





👤 Controls

👤 Controls

🕒 Kalshi prior (predictions)

🕒 Kalshi prior (predictions)

🌡️ Temperature 0.2

🌡️ Temperature 0.2

Home Assistant, over MQTT

Three accounts

official, test and mixed — one section each, the same entities

Kill switch

flatten and halt, **that account only**; all three stays CLI-only

State

equity, day P&L, open positions, halt state, last decision

Day P&L

a history graph, per account

Control knobs

stop-loss, take-profit, EOD close DTE, min hold, early-exit drawdown, strike band, research breadth, temperature, Kalshi prior — live, expiring at the close

AI controls

model and review model, max tokens, tokens today and all time

Live journal

the cycle feed as it lands

Trades & EOD

today's fills and rejections, then the digest

On a phone

the grid reflows to one account per screen; halts, identity refusals and 25 minutes of *silence* push — fills do not


Remote access

the dashboard lives on the home LAN; judges can have it **on request** — `ha.wpmccormick.pw` through the same Cloudflare Tunnel, Cloudflare Access allow-listed by email address, no one-time PIN

Autobelay long premium, short leash

page 013

Our own hardware, not a laptop



Built to be changed while it runs

795 tests, 2s

every one runs with **no API keys and no network**

Only Docker needed

tests, a full dry-run cycle, the docs site

Rehearse any time

`--dry-run` prints orders instead of sending; `--force` skips the market gates

No deploy needed

`override.py`: allowlisted knobs, hard caps stay git-only, and all of it expires at the close

Dozens of those tests exist because something broke. Each incident ends as a test.

The daily loop



The critique comes from a model that did *not* trade the day. That loop ran 180 pull requests and 132 deploys in six days.

Choosing the model, honestly

Featherless lists **21,912** models, **15,575** with tool calling on our plan. "Open weights" narrows nothing; our gates do, and none of them is a benchmark score.

● THE GATES

Tool calling for the research loop · context for a 5k prompt that reaches **95k** tokens mid-research · a thinking-mode toggle, without which these models answer *empty* · a warm endpoint for a 10-minute cron · a price, so the day's spend is reported.

`scripts/verify_models.py` re-checks every one against the live catalog, so the claim cannot rot.

● WHAT WE WILL NOT CLAIM


No bake-off picked the incumbent. Kimi was the first model confirmed to support tool calling, and it stayed. Our A/B varies the model *and* the tools *and* the learning loop, so it compares config bundles.

Saying so beats inventing a rationale for it.

The gate that discriminates is **instruction adherence**. `Llama-3.1-Hawkish-8B` — a finance fine-tune, cheaper than everything we run, clears every mechanical gate — proposed a 1-DTE contract in **3 of 3** runs and five actions against a four-position cap. Well-formed JSON, wrong content. It also found a real bug: the funnel permitted what the prompt forbade.

Real A/B, not a backtest

Four days is no backtest, and the free feed carries no historical options data, so the experiment runs live instead.




same market, same cadence, one difference each

● THREE LIVE ACCOUNTS

Same box, same ten-minute cadence, same market, one deliberate difference each. `test` swaps the model and enables research and learning; `mixed` differs by exactly one key.

The winner is promoted into the official config by pull request.

Results



Judged account Mon open → Thu close: **\$100,000 → \$105,668 (+5.7%) · 15**
round trips · what didn't work: the Kimi-K3 challenger (−7.1%, hit the 2%
cutoff Thu) and exit-then-re-enter churn on QQQ and NVDA

What's next

- **PROMOTE WHAT THE CHALLENGERS PROVED**

Research tools, the prediction-market prior, the learning loop.

- **MULTI-LEG SPREADS**

Once the gates cover assignment risk.

github.com/bill-mccormick-dg/alpaca-hackathon · MIT

Thanks to Alpaca, lablab.ai and Featherless AI.